

Walmart

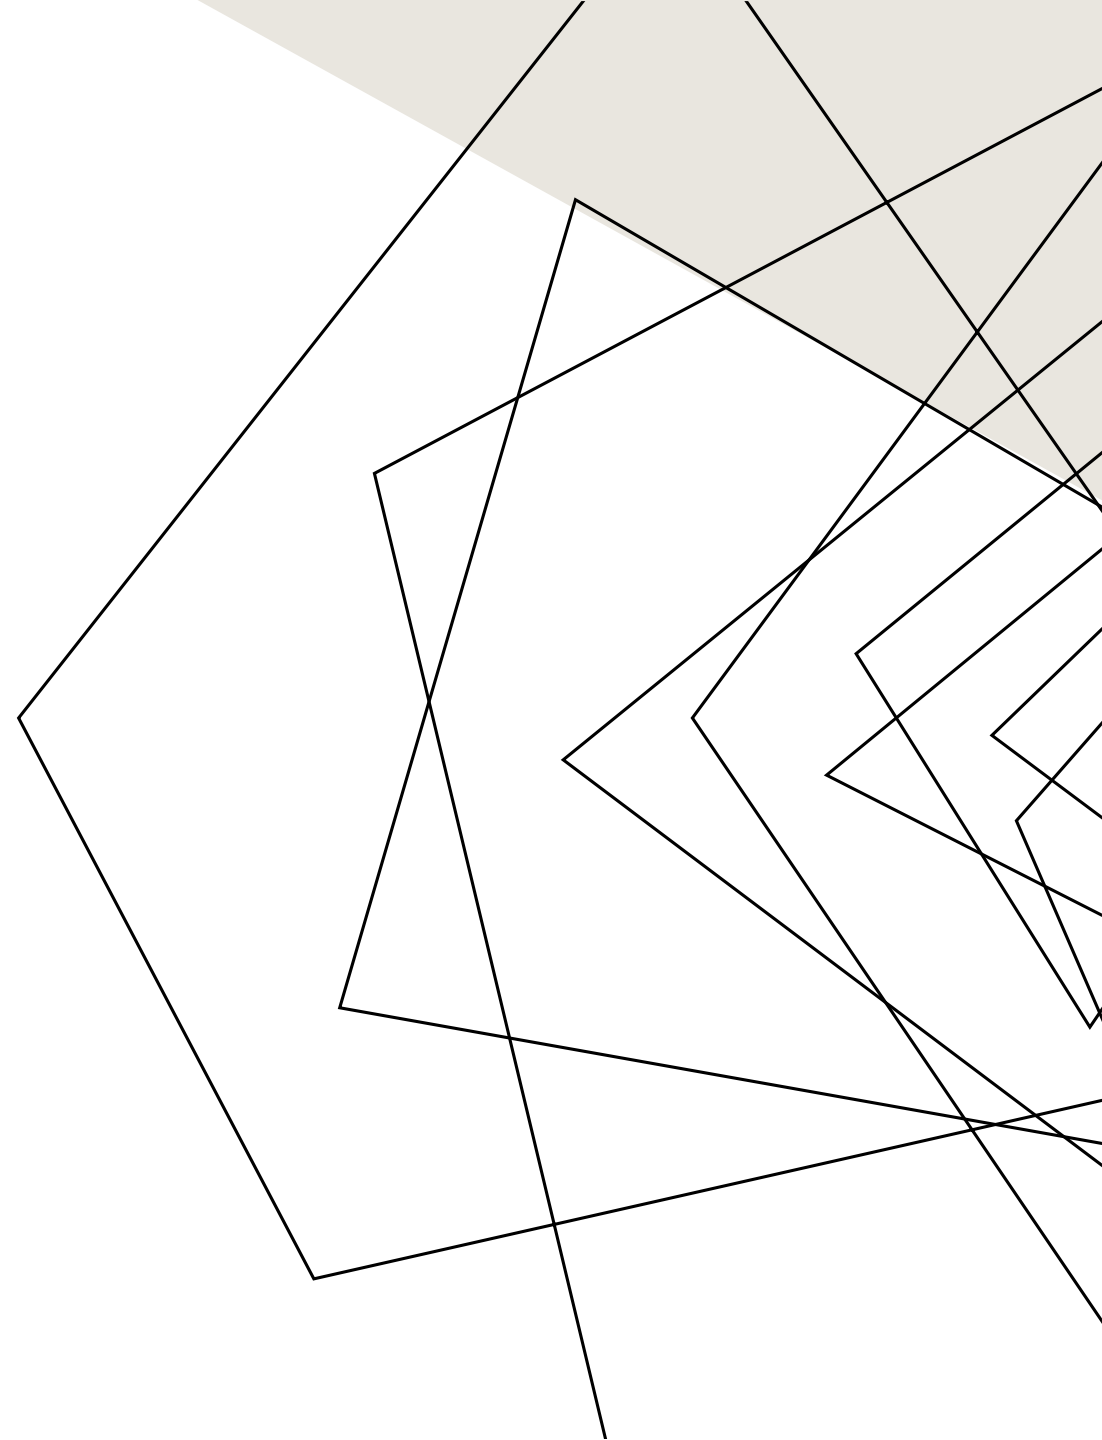
The Walmart logo is displayed in large, white, three-dimensional letters on a blue background. To the right of the word "Walmart" is the Walmart spark logo, which consists of eight yellow, three-dimensional, chevron-like shapes arranged in a circular pattern.

5250 Commercial St. SE



WALMART

We decided to predict the differences in weekly sales of Walmart the american retailer



TABLES OF THE SALES AND THE STORES INFORMATION

```
[18] store_df.head()
```

	Store	Type	Size
0	1	A	151315
1	2	A	202307
2	3	B	37392
3	4	A	205863
4	5	B	34875



```
df.head()
```



	Store	Dept	Date	Weekly_Sales	IsHoliday
0	1	1	2010-02-05	24924.50	False
1	1	1	2010-02-12	46039.49	True
2	1	1	2010-02-19	41595.55	False
3	1	1	2010-02-26	19403.54	False
4	1	1	2010-03-05	21827.90	False

ANOTHER TABLE THAT HAS INFORMATION ON THE STORE IN A SPECIFIC WEEK

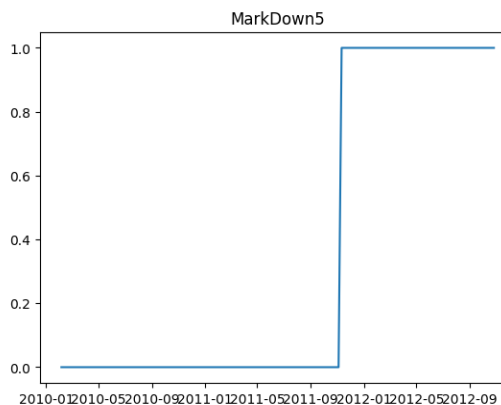


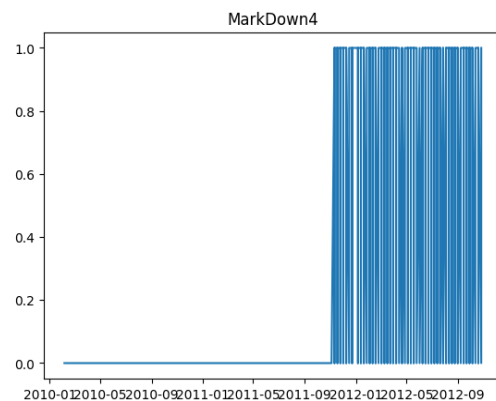

```
[17] features_df.head()
```

	Store	Date	Temperature	Fuel_Price	MarkDown1	MarkDown2	MarkDown3	MarkDown4	MarkDown5	CPI	Unemployment	IsHoliday
0	1	2010-02-05	42.31	2.572	NaN	NaN	NaN	NaN	NaN	211.096358	8.106	False
1	1	2010-02-12	38.51	2.548	NaN	NaN	NaN	NaN	NaN	211.242170	8.106	True
2	1	2010-02-19	39.93	2.514	NaN	NaN	NaN	NaN	NaN	211.289143	8.106	False
3	1	2010-02-26	46.63	2.561	NaN	NaN	NaN	NaN	NaN	211.319643	8.106	False
4	1	2010-03-05	46.50	2.625	NaN	NaN	NaN	NaN	NaN	211.350143	8.106	False

We decided to remove features

- Markdown1
- Markdown2
- Markdown3
- Markdown4
- Markdown5

Since they exist only since 2012 and even then
are lacking



JOINED TABLE

	Store	Dept	Date	Weekly_Sales	IsHoliday_x	Type	Size	Temperature	Fuel_Price	CPI	Unemployment	IsHoliday_y
0	1	1	2010-02-05	24924.50	False	A	151315	42.31	2.572	211.096358	8.106	False
1	1	2	2010-02-05	50605.27	False	A	151315	42.31	2.572	211.096358	8.106	False
2	1	3	2010-02-05	13740.12	False	A	151315	42.31	2.572	211.096358	8.106	False
3	1	4	2010-02-05	39954.04	False	A	151315	42.31	2.572	211.096358	8.106	False
4	1	5	2010-02-05	32229.38	False	A	151315	42.31	2.572	211.096358	8.106	False
...
421565	45	93	2012-10-26	2487.80	False	B	118221	58.85	3.882	192.308899	8.667	False
421566	45	94	2012-10-26	5203.31	False	B	118221	58.85	3.882	192.308899	8.667	False
421567	45	95	2012-10-26	56017.47	False	B	118221	58.85	3.882	192.308899	8.667	False
421568	45	97	2012-10-26	6817.48	False	B	118221	58.85	3.882	192.308899	8.667	False
421569	45	98	2012-10-26	1076.80	False	B	118221	58.85	3.882	192.308899	8.667	False

421570 rows × 12 columns


GENERAL STATISTIC INFORMATION

[24] df.describe()


	Store	Dept	Weekly_Sales	Size	Temperature	Fuel_Price	CPI	Unemployment
count	421570.000000	421570.000000	421570.000000	421570.000000	421570.000000	421570.000000	421570.000000	421570.000000
mean	22.200546	44.260317	15981.258123	136727.915739	60.090059	3.361027	171.201947	7.960289
std	12.785297	30.492054	22711.183519	60980.583328	18.447931	0.458515	39.159276	1.863296
min	1.000000	1.000000	-4988.940000	34875.000000	-2.060000	2.472000	126.064000	3.879000
25%	11.000000	18.000000	2079.650000	93638.000000	46.680000	2.933000	132.022667	6.891000
50%	22.000000	37.000000	7612.030000	140167.000000	62.090000	3.452000	182.318780	7.866000
75%	33.000000	74.000000	20205.852500	202505.000000	74.280000	3.738000	212.416993	8.572000
max	45.000000	99.000000	693099.360000	219622.000000	100.140000	4.468000	227.232807	14.313000

KEEPING TRACK OF TIME WITH SIN AND COS

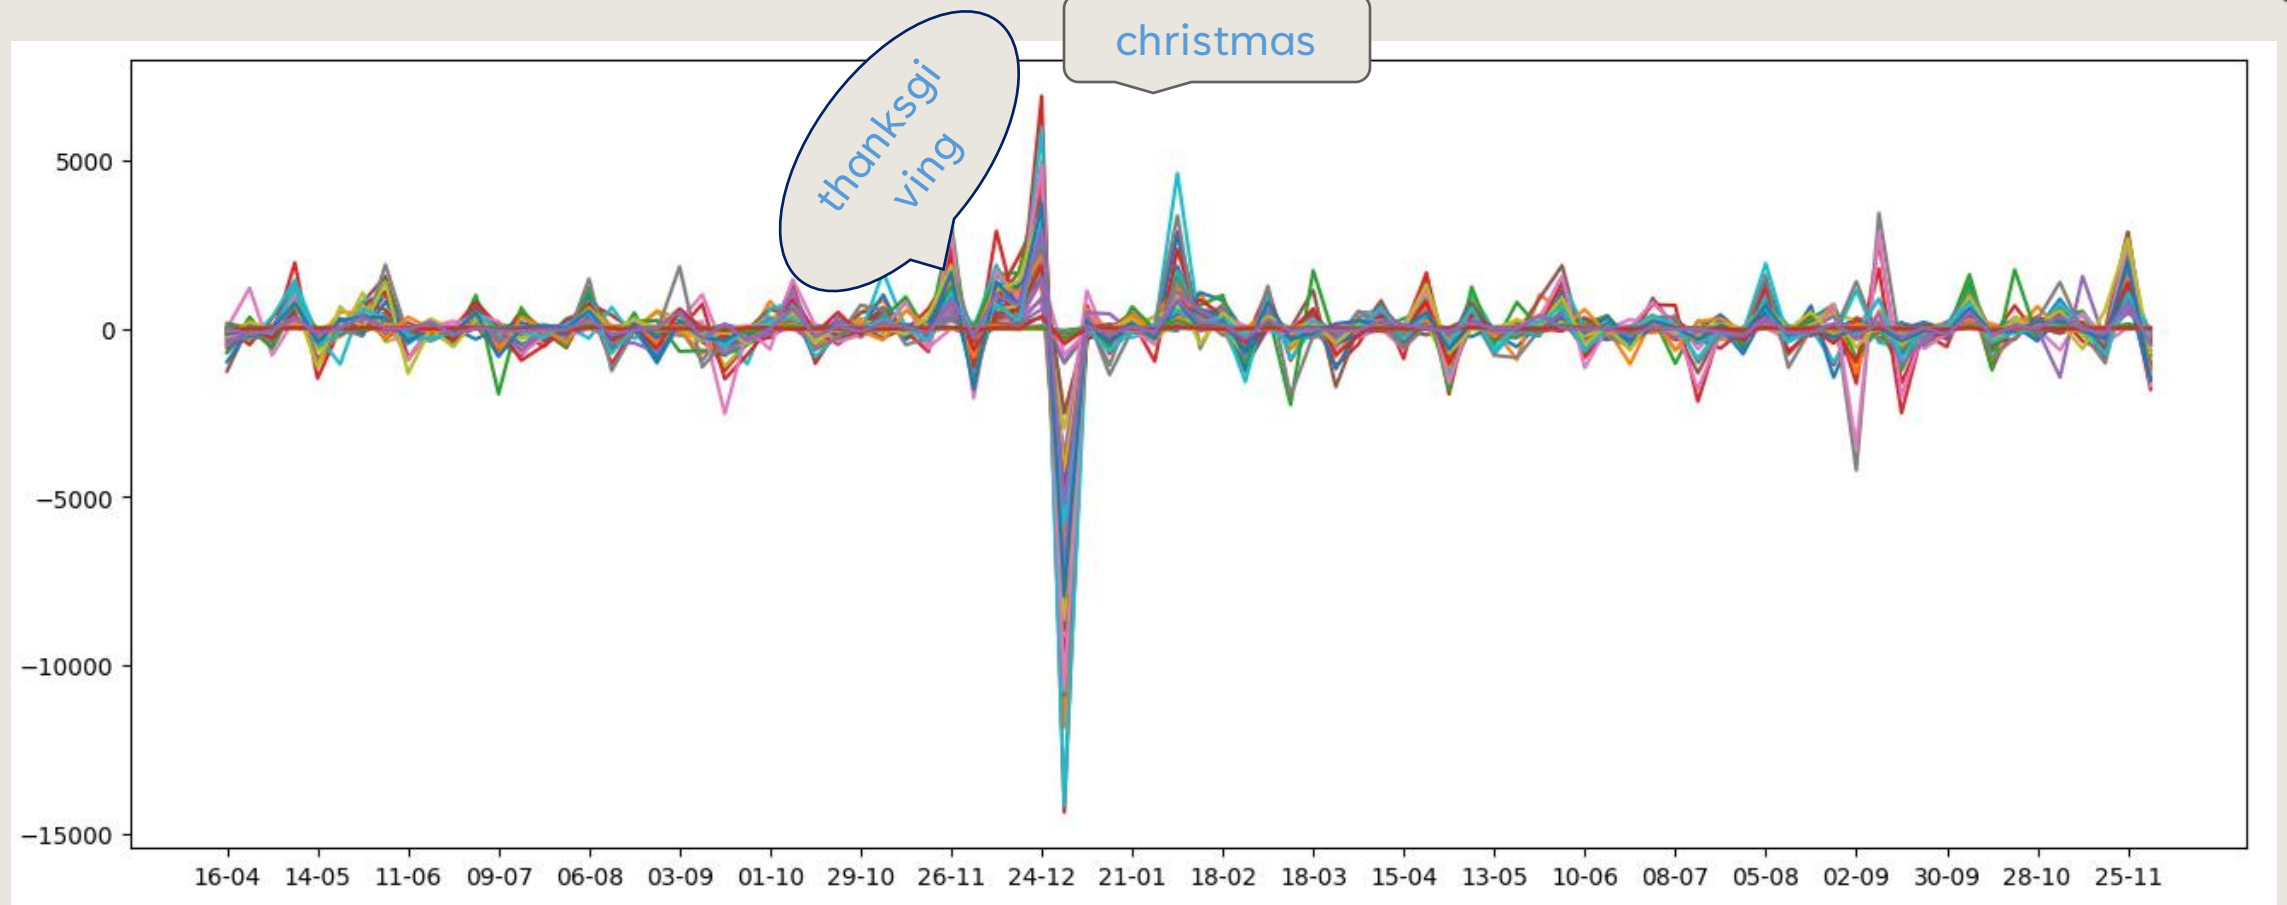
```
plt.figure(figsize=(16, 6))  
  
plt.plot(train_df["Date"], train_df["Year-Sin"], '.')  
plt.plot(train_df["Date"], train_df["Year-Cos"], '.')  
  
plt.show()
```







MEDIAN OF THE SALES FOR EACH DEPARTMENT







MEDIAN OF THE WEEKLY DIFFERENCE IN SALES FOR EACH DEPARTMENT



LOOKING AT SALES OF ALL DEPARTMENTS FOR DIFFERENT STORES



LOOKING AT ALL STORES FOR DIFFERENT DEPARTMENTS



Linear Regression data

index	Store	Dept	Date	Temperature	Fuel_Price	CPI	Unemployment	IsHoliday	Store-Type	Store-Size
0	1	1	2010-05-28	80.44	2.759	210.896761	7.808	0	1	0.233352
1	1	1	2010-06-04	80.69	2.705	211.176428	7.808	0	1	0.233352
2	1	1	2010-06-11	80.43	2.668	211.456095	7.808	0	1	0.233352
3	1	1	2010-06-18	84.11	2.637	211.453772	7.808	0	1	0.233352
4	1	1	2010-06-25	84.34	2.653	211.338653	7.808	0	1	0.233352

Weekly-Sales-last-1-weeks	Weekly-Sales-last-2-weeks	Weekly-Sales-last-3-weeks	Weekly-Sales-last-4-weeks	Weekly-Diff-last-1-weeks	Weekly-Diff-last-2-weeks	Weekly-Diff-last-3-weeks	Weekly-Diff
14773.04	18926.74	17413.94	16555.11	-4153.70	1512.80	858.83	807.39
15580.43	14773.04	18926.74	17413.94	807.39	-4153.70	1512.80	1977.66
17558.09	15580.43	14773.04	18926.74	1977.66	807.39	-4153.70	-920.47
16637.62	17558.09	15580.43	14773.04	-920.47	1977.66	807.39	-421.35
16216.27	16637.62	17558.09	15580.43	-421.35	-920.47	1977.66	112.45

We will use Stochastic Gradient
Descent in order to learn linear
regression

We will train with MSE as
out loss function but use
RMSE to evaluate

model

```
def train_linear_regression_model(x_train, y_train, x_val, y_val, model, loss_fn, loss_fn_val, num_epochs, batch_size):
    lr = 0.01

    # Define an optimizer (Stochastic Gradient Descent)
    optimizer = torch.optim.SGD(model.parameters(), lr=lr)

    train_dataset = torch.utils.data.TensorDataset(torch.tensor(x_train, dtype=torch.float32),
                                                    torch.tensor(y_train, dtype=torch.float32))
    train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)

    val_dataset = torch.utils.data.TensorDataset(torch.tensor(x_val, dtype=torch.float32),
                                                  torch.tensor(y_val, dtype=torch.float32))
    val_loader = torch.utils.data.DataLoader(val_dataset, batch_size=batch_size, shuffle=True)

    losses = []
    losses_val = []

    for epoch in range(num_epochs):
        running_loss = 0

        for batch in train_loader:
            inputs, targets = batch

            # Forward pass
            outputs = model(inputs)
            loss_train = loss_fn(outputs, targets)

            # Backward and optimize
            optimizer.zero_grad()
            loss_train.backward()
            optimizer.step()

            running_loss += np.sqrt(loss_train.item())

        losses.append(running_loss / len(train_loader))

        # Validation loss
        running_loss = 0

        for batch in val_loader:
            inputs, targets = batch

            # Forward pass
            outputs = model(inputs)
            loss_val = loss_fn_val(outputs, targets)

            running_loss += np.sqrt(loss_val.item())

        losses_val.append(running_loss / len(val_loader))

        if epoch % int(num_epochs / 10) == 0:
            print(f'Epoch [{epoch}], Running Loss: {running_loss:.4f}')

    return losses, losses_val
```

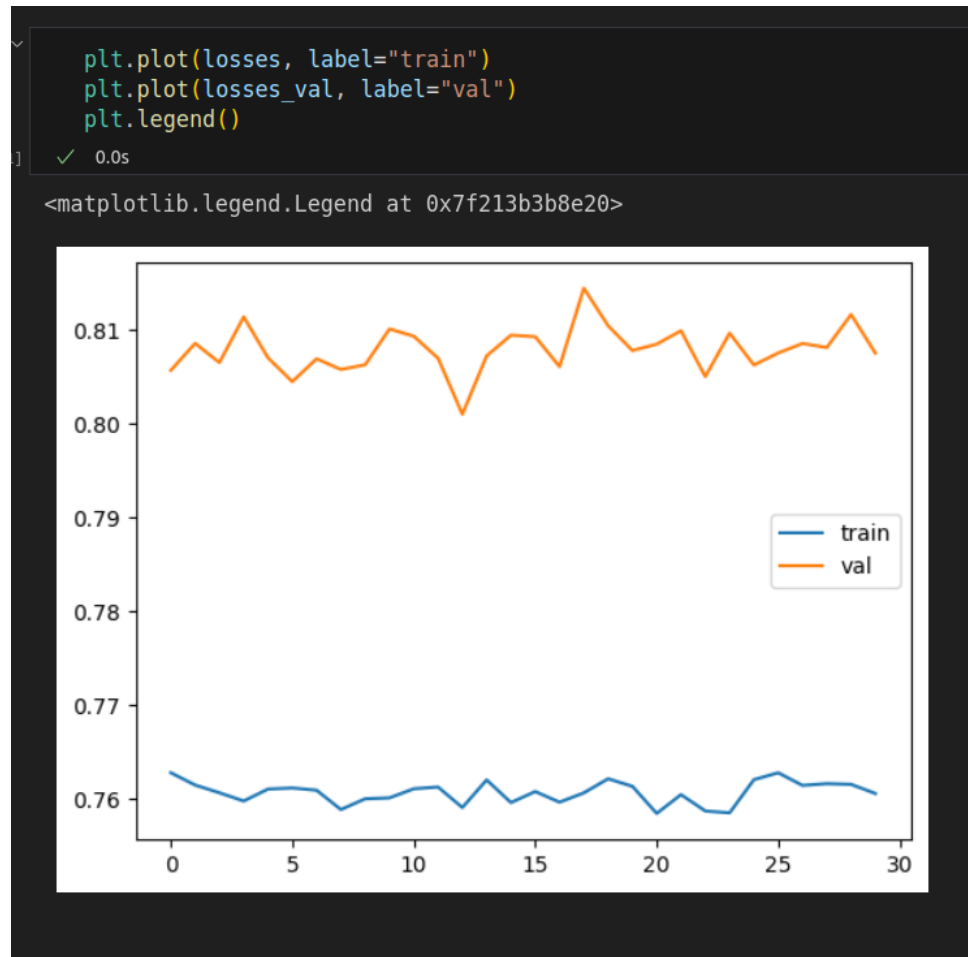
```
class LinearRegression(torch.nn.Module):
    def __init__(self, input_size):
        super(LinearRegression, self).__init__()
        self.linear = torch.nn.Linear(input_size, 1)

    def forward(self, x):
        out = self.linear(x)
        return out
```


training

Loss function


Is it really
learning
?



normalized



Non normalized



```

class NormlizerStdMean:
    def __init__(self, x, y):
        self.y_mean = y.mean()
        self.y_std = y.std()

        self.x_mean = {}
        self.x_std = {}

        for column in x.columns:
            self.x_mean[column] = x[column].mean()
            self.x_std[column] = x[column].std()

    def normalize_y(self, v):
        return (v - self.y_mean) / self.y_std

    def normalize_x(self, v, column):
        return (v - self.x_mean[column]) / self.x_std[column]

    def denormalize_y(self, v):
        return list(map(lambda v0: (v0 * self.y_std) + self.y_mean, v))

    def denormalize_x(self, v, column):
        return list(map(lambda v0: (v0 * self.x_std[column]) + self.x_mean[column], v))

```

```

class NormlizerMinMax:
    def __init__(self, x, y):
        self.y_min = y.min()
        self.y_max = y.max()

        self.x_min = {}
        self.x_max = {}

        for column in x.columns:
            self.x_min[column] = x[column].min()
            self.x_max[column] = x[column].max()

    def normalize_y(self, v):
        return (v - self.y_min) / (self.y_max - self.y_min)

    def normalize_x(self, v, column):
        return (v - self.x_min[column]) / (self.x_max[column] - self.x_min[column])

    def denormalize_y(self, v):
        return list(map(lambda v0: (v0 * (self.y_max - self.y_min)) + self.y_min, v))

    def denormalize_x(self, v, column):
        return list(map(lambda v0: (v0 * (self.x_max[column] - self.x_min[column])) + self.x_min[column], v))

```

```

def train_linear_regression_model(x_train, y_train, x_val, y_val, model, loss_fn, loss_fn_val, num_epochs, batch_size):
    lr = 0.01

    # Define an optimizer (Stochastic Gradient Descent)
    optimizer = torch.optim.SGD(model.parameters(), lr=lr)

    train_dataset = torch.utils.data.TensorDataset(torch.tensor(x_train, dtype=torch.float32),
                                                    torch.tensor(y_train, dtype=torch.float32))
    train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)

    val_dataset = torch.utils.data.TensorDataset(torch.tensor(x_val, dtype=torch.float32),
                                                  torch.tensor(y_val, dtype=torch.float32))
    val_loader = torch.utils.data.DataLoader(val_dataset, batch_size=batch_size, shuffle=True)

    losses = []
    losses_val = []

    for epoch in range(num_epochs):
        running_loss = 0

        for batch in train_loader:
            inputs, targets = batch

            # Forward pass
            outputs = model(inputs)
            loss_train = loss_fn(outputs, targets)

            # Backward and optimize
            optimizer.zero_grad()
            loss_train.backward()
            optimizer.step()

            running_loss += np.sqrt(loss_train.item())

        losses.append(running_loss / len(train_loader))

        # Validation loss
        running_loss = 0

        for batch in val_loader:
            inputs, targets = batch

            # Forward pass
            outputs = model(inputs)
            loss_val = loss_fn_val(outputs, targets)

            running_loss += np.sqrt(loss_val.item())

        losses_val.append(running_loss / len(val_loader))

        if epoch % int(num_epochs / 10) == 0:
            print(f'Epoch [{epoch}], Running Loss: {running_loss:.4f}')

    return losses, losses_val

```

Model(same as before)

```

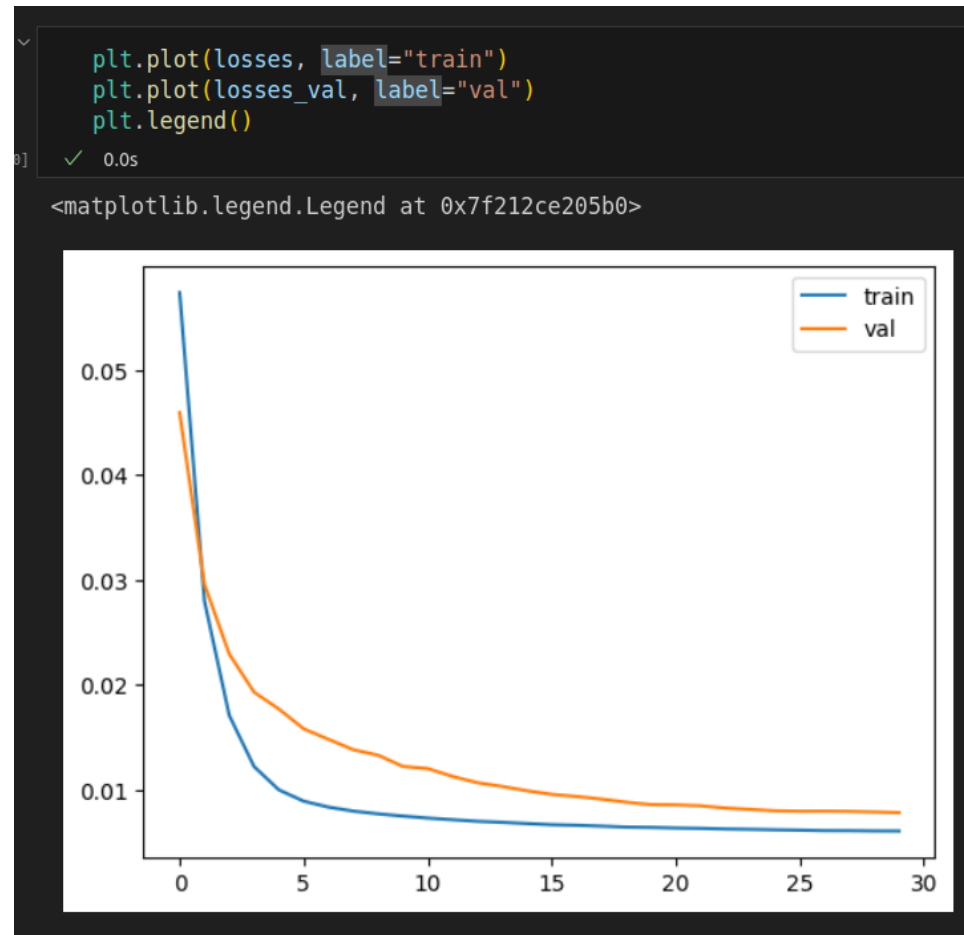
class LinearRegression(torch.nn.Module):
    def __init__(self, input_size):
        super(LinearRegression, self).__init__()
        self.linear = torch.nn.Linear(input_size, 1)

    def forward(self, x):
        out = self.linear(x)
        return out


```

Training also the same


Loss function before and after test





normalized



Non
normalized



TIME PREDICTION IS NOT EASY

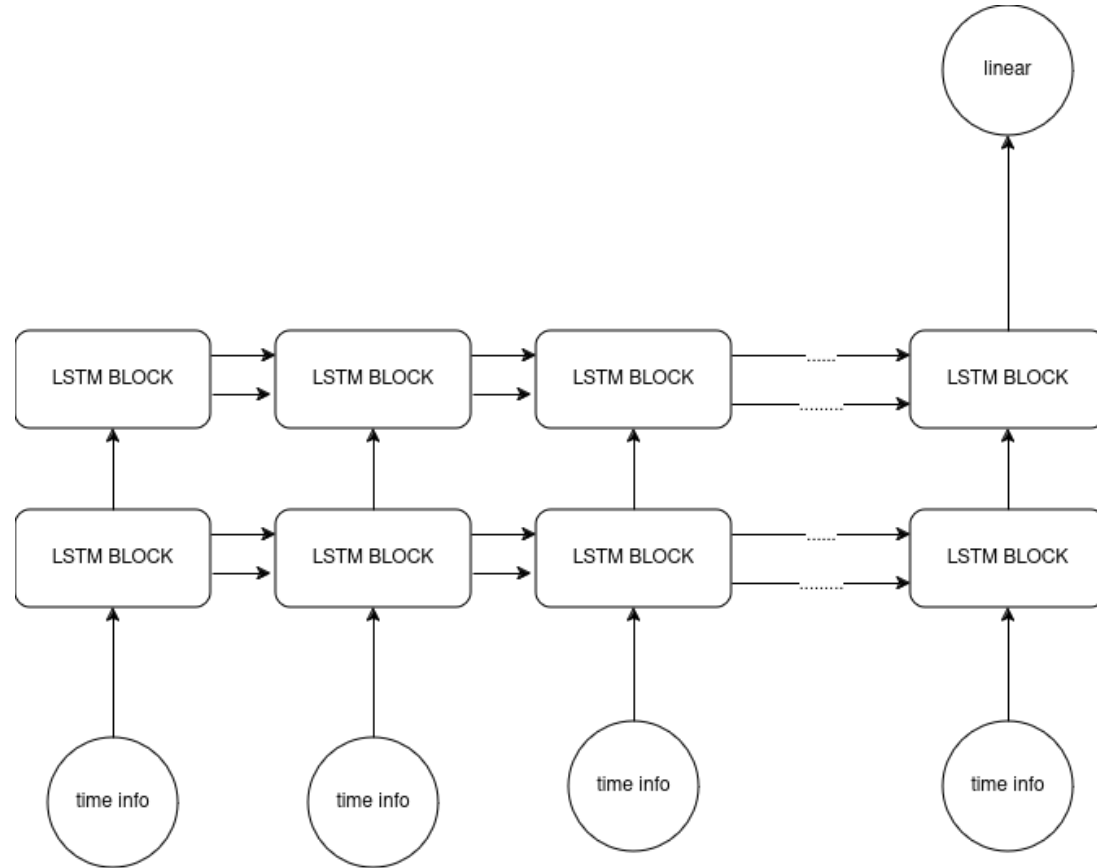


RNN(SPECIFICALLY LSTM) TO THE RESCUE

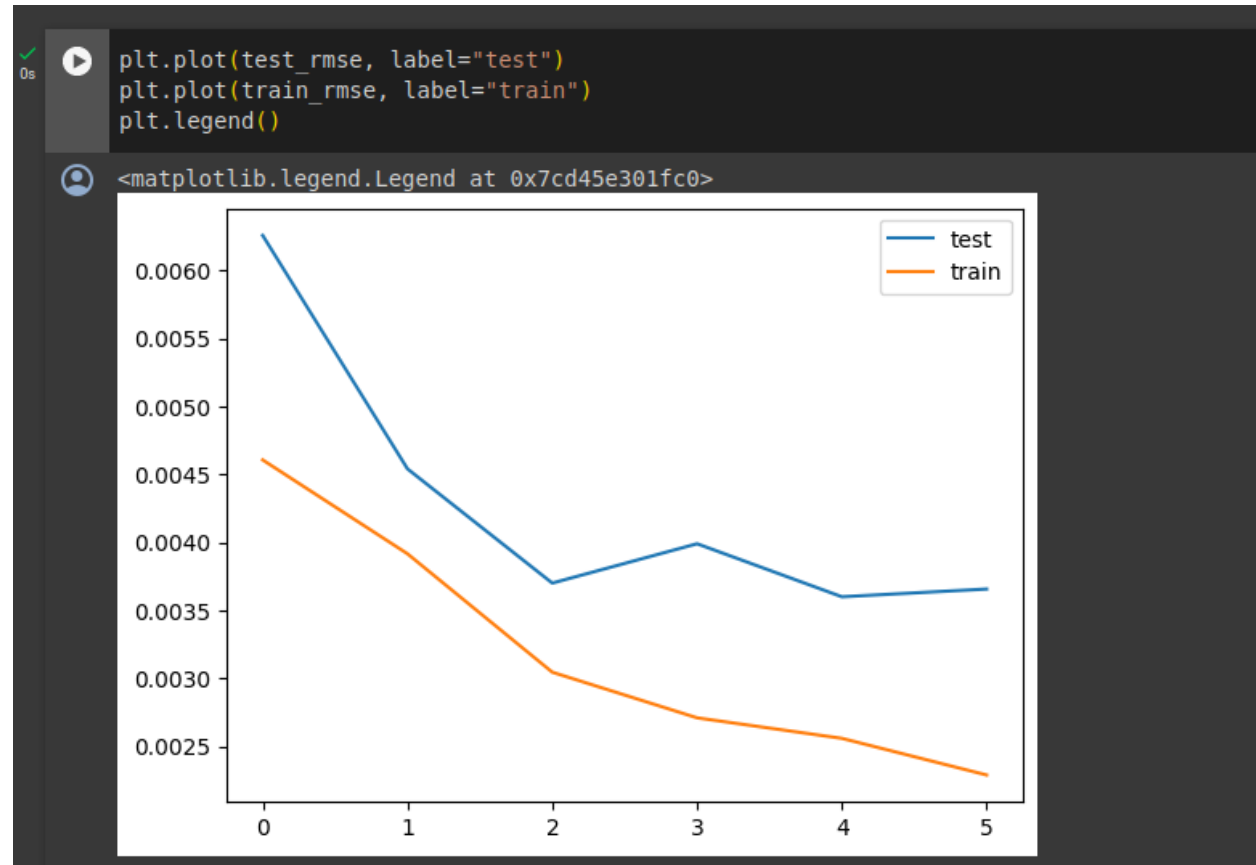
- Has built in time flow
- Is proven to be effective
- Already solved the vanishing gradient problem
- Has comparatively low number of weights for the amount of data

OUR ARCHITECTURE


- Two layered LSTM
- Hidden state passed through each block =50
- At each time period we get the all the relveant data from today(including sin cos of the date instead of the date)
- Additionally we add the last week's week-diff



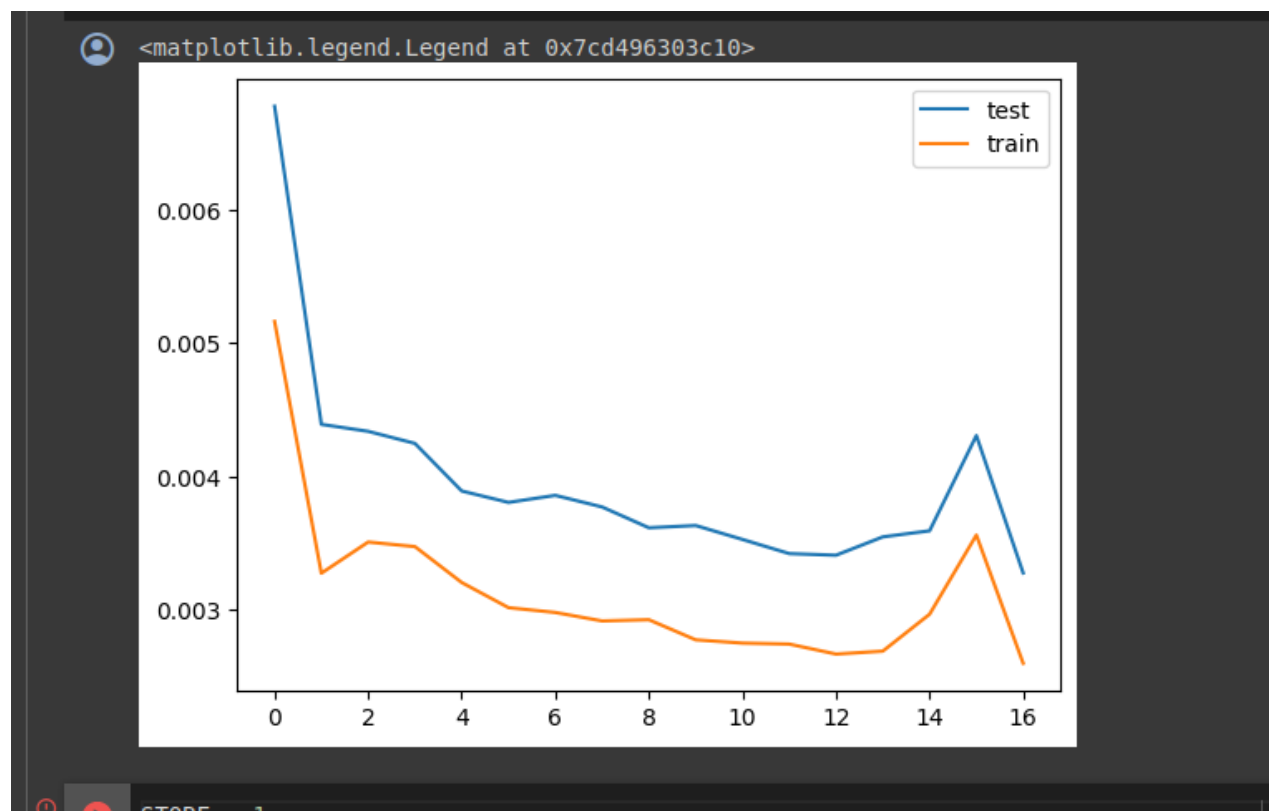
RESULT AFTER 5 EPOCHS WITH 5 LSTM LAYERS (SEEMS TO BE OVERFITTING BUT CANT BE SURE)




RESULTS AFTER 30 EPOCHS WITH ONLY 2 LSTM LAYERS(OBSERVED EVERY 3 EPOCHS)





NEED TO SOLVE OVERFITTING PROBLEM....
DROPOUT!!!(0.2) AND 50 EPOCHS(OBSERVED EVERY 3 EPOCHS)





OBSERVED RESULTS IN ARBITRARY STORE AND SPECIFIC DEPARTMENT (DEPT=3)




OBSERVED RESULTS IN THE SAME DEPARTMENT BUT AT A DIFFERENT STORE




RESULTS IN A DIFFERENT DEPARTMENT




USING LOOKBACK = 10 AND 150 EPOCHS



RESULTS



RESULTS





AND NOW CLAP FOR THE
FINAL SCORE

An abstract geometric pattern consisting of several white lines of varying lengths and orientations, creating a complex, overlapping structure on the left side of the image. The lines form various polygons and intersect to create a sense of depth and movement.

THANK YOU

We got Loss of 0.0026

Lior Shiboli and Omer Priel